

elastic Meetup Night

Class: Centralized Logging

Motivation



- Aggregate log events
 - across applications
 - across processes
 - across machines
- Combine and visualize data

Use Cases



- Business
 - amount and duration of transactions
 - history of operations
 - correlation between applications
- Performance and health
 - amount and duration of requests, request times
 - errors
 - correlation between systems, data bases, applications
- Security
 - fraud detection, intrusion detection
 - tracing

Overview



Beats



- Lightweight agent that pushes data to another destination
- Common for centralized logging: Filebeat
- Implemented in Go
- Other beats:
 - Metricbeat
 - Winlogbeat
 - Packetbeat

Logstash



- Kind of ETL system for transforming event data
- Often used for logs
- Written in JRuby
- Configured in Ruby Syntax
 - Input
 - Filter
 - Output

Elasticsearch



- Stores the log events
- Query DSL used for accessing the data
- Often time based indexing
 - filebeat-2017-03-23
 - Index pattern filebeat-*

Kibana



- Web interface for analyzing data
- Explore data
- Generates requests to Elasticsearch
- Visualizes results: diagrams, maps
- Written in JavaScript

Ingest Node



- Logstash can consume quite some resources
- Common transformations available in Elasticsearch as well
- Ingestion can happen on separate node

Reading files: Filebeat



- Installation as service or standalone
- Configuration in **filebeat.yml**
- Prospectors define where to read data from
- Output determines where to send data to



Filebeat: Prospector configuration

```
filebeat.prospectors:
```

```
- input_type: log
```

```
  paths:
```

```
    - /var/log/*.log
```

```
    #- c:\programdata\elasticsearch\logs\*
```


- Includes, excludes
- Multiline configuration

Filebeat: Setup options

- Using Ingest Node to process data



- Using Logstash to process data





Filebeat: Output configuration

- Sending data to Elasticsearch
 - for storing the message as a whole
 - or for processing with Ingest pipeline

```
output.elasticsearch:  
  hosts: ["localhost:9200"]  
  
  # Optional protocol and basic auth credentials.  
  # username: "elastic"  
  # password: "changeme"
```

Filebeat: Output configuration



- Sending data to Logstash
 - for processing and redistribution

```
output.logstash:  
  # The Logstash hosts  
  hosts: ["localhost:5044"]
```



Reading files: Logstash

- 3 major sections
 - Input: Specifies where to read data from
 - Filter: Different ways to massage the data
 - Output: Data sinks
- Countless plugins for all

Logstash: Input



- Server process for reading beat data

```
input {  
  beats {  
    port => 5044  
  }  
}
```

- or reading from file

```
input {  
  file {  
    path => "/var/log/apache2/access.log"  
  }  
}
```


Logstash: Grok filter



- Common filter: Grok
 - Regular expressions with predefined patterns

```
filter {  
  grok {  
    match => { message => "%{COMBINEDAPACHELOG}" }  
  }  
}
```

Logstash: Output



- Sending data to Elasticsearch

```
output {  
  elasticsearch {  
    host => "localhost"  
    protocol => "http"  
  }  
}
```

- For debugging

```
stdout {  
  codec => rubydebug  
}
```

Ingest Node



- Implementation of some of the Logstash filters
- Pipeline defines processing order
- management via HTTP API
- Pipeline can be defined in Filebeat config

Ingest Node



- Pipeline has name and processors
 - simple case: number followed by a word

```
PUT _ingest/pipeline/simplest-log
{
  "processors" : [
    {
      "grok" : {
        "field": "message",
        "patterns": ["%{NUMBER:line} %{WORD:text}"]
      }
    }
  ]
}
```

- Result: documents with line and text field

Ingest Node



- Configure pipeline in Filebeat

```
output.elasticsearch:  
  hosts: ["localhost:9200"]  
  pipeline: simplest-log
```

Ingest Node



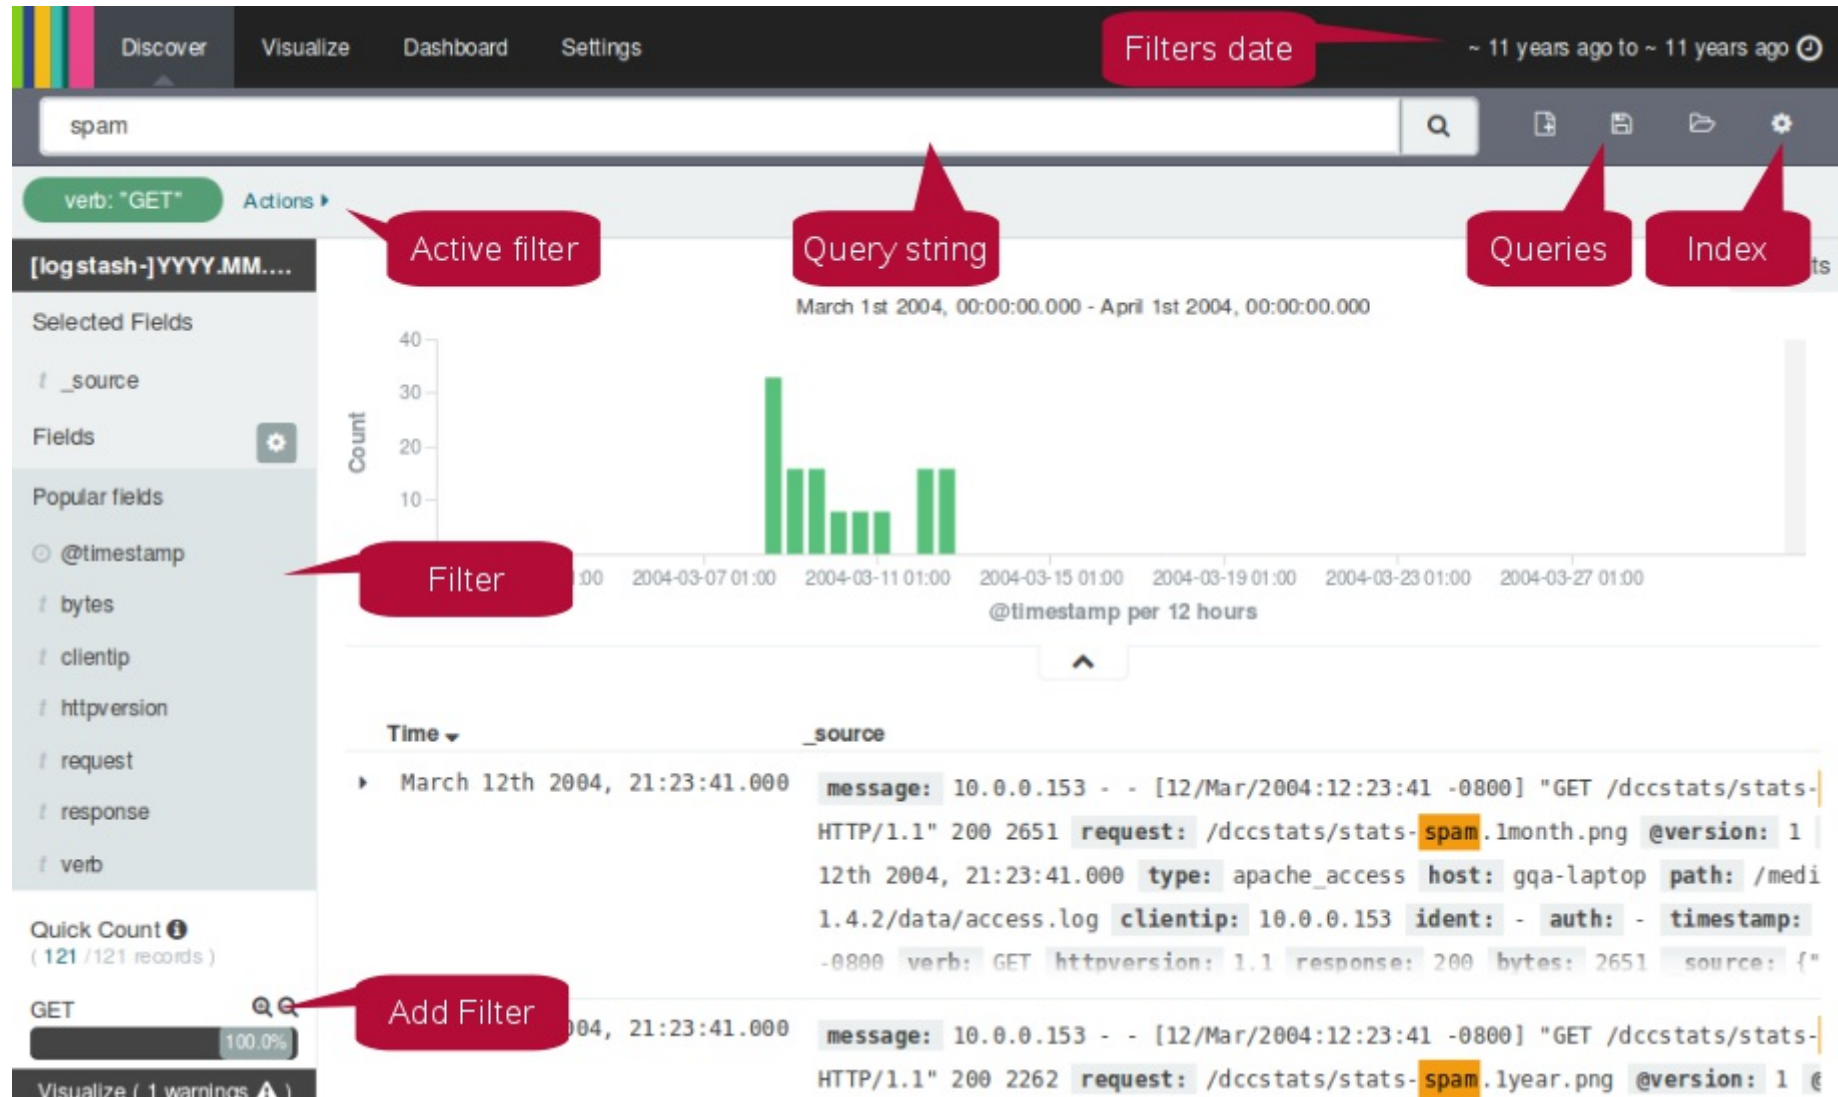
- Simulation for debugging
- Online tools for creating grok expressions
 - grokdebug <https://grokdebug.herokuapp.com/>
 - Grok Constructor <http://grokconstructor.appspot.com/>
- Careful with expensive regular expressions

Kibana

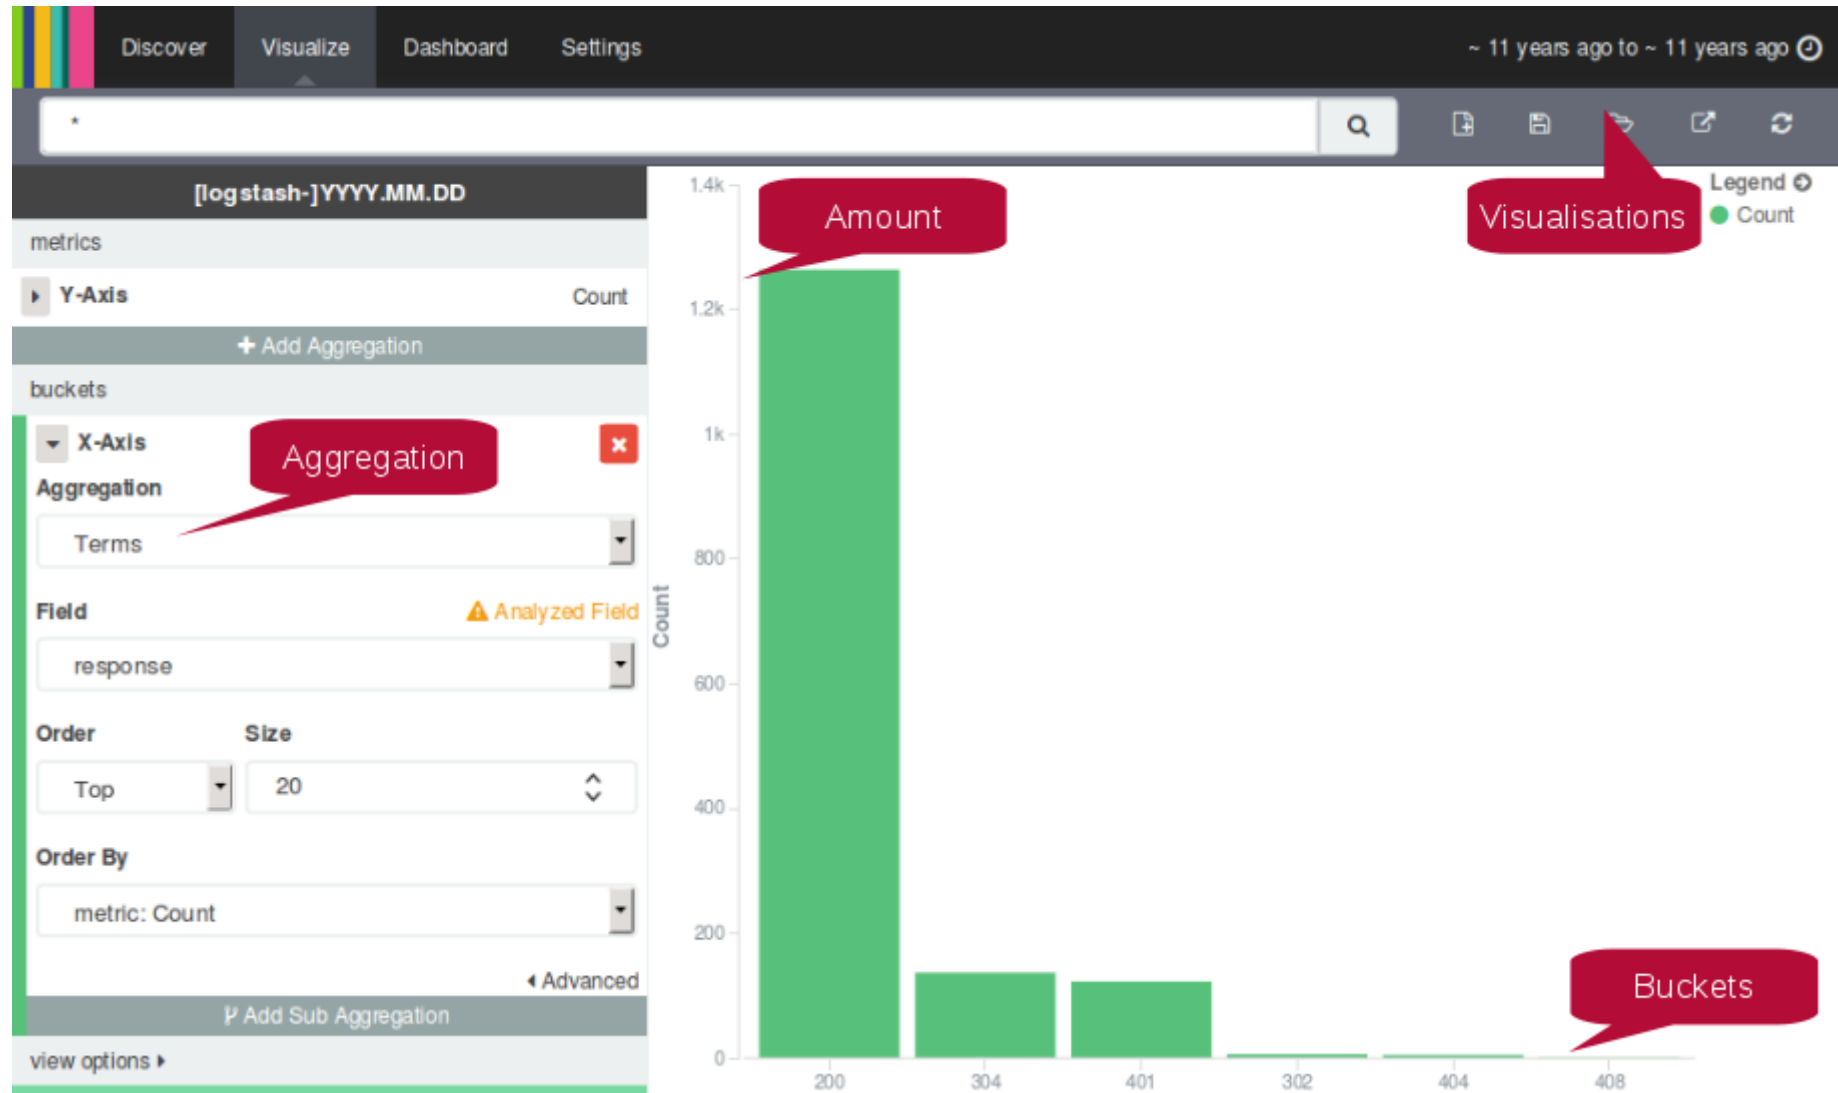


- Search in log events
- Visualize data

Kibana



Kibana



Scaling options



- Backpressure sensitive when sending to Logstash and Elasticsearch
- Buffer for load spikes and for scaling Logstash
 - Redis
 - Kafka

Scaling options

